

ANIMA

A Deterministic Behavioral Organization Engine
for Structurally Consistent LLM Agents

Technical Report · Project ANIMA · 2026

Confidential — Early-stage research · Not for redistribution

Abstract

We present ANIMA (Generative Motor of Structured Subjectivity), an architecture that decouples behavioral organization from language generation in LLM-based agents. In current agent frameworks, behavioral consistency is maintained stochastically by the language model itself — through prompt engineering and context windows — with no persistent internal state, no accumulating pressure, and no structural history. ANIMA introduces a deterministic psychodynamic engine that computes a seven-dimensional state vector $S(t) \in [0,1]^7$ at each conversational turn, then instructs the LLM to render that state as natural language. The LLM is a renderer; ANIMA is the agent.

The architecture comprises three components: (1) a psychodynamic engine with deterministic update equations governing seven state variables; (2) an LLM coupling protocol that translates engine state into precise structured instructions; and (3) a Deterministic Signal Extractor (DSE) that computes discourse signals from rendered text without additional LLM calls. We demonstrate the architecture across seven behavioral archetypes derived from psychoanalytic structural theory and present a blind clinical evaluation protocol designed to measure structural distinguishability. Applications include agent behavioral evaluation, structured synthetic data generation, and persistent NPC design.

1. Introduction: The Behavioral Consistency Problem

Large language models have transformed conversational AI. Yet their deployment as long-running agents reveals a fundamental limitation: behavioral consistency is not a property of the model — it is an artifact of the context window.

When a system prompt instructs a model to behave as a particular type of agent, three structural problems emerge independently of model capability:

- No accumulating internal state. The model cannot 'need' to say something across turns. It cannot build pressure that eventually breaks through. There is no mechanism for content that exists but cannot yet be expressed.
- No structural history. The agent in turn 50 is not systematically different from the agent in turn 1. The model has no memory of its own trajectory — only of the conversational record.
- Non-reproducibility. The same state, the same question, different behavioral trajectories. Behavioral evaluation is undermined by stochasticity at the generation layer.

These are not engineering failures. They are consequences of the architecture: language models are stateless renderers. Asking a renderer to maintain structural consistency over time is asking it to do something it is not built to do.

ANIMA's intervention is architectural. Rather than asking the LLM to maintain behavioral state, ANIMA computes that state externally through a deterministic engine and instructs the LLM to express it. The LLM does what it does best — produce fluent natural language from instructions. The engine does something the LLM cannot: maintain a persistent, evolving state with internal pressure, structural history, and reproducible dynamics.

The theoretical framework that informed ANIMA's engine design is psychoanalytic: specifically, the structural theory of clinical types developed by Freud, Lacan, and successors. We chose this framework not for disciplinary reasons but because it offers the most developed formal account of how internal states organize discourse — how a subject's internal economy determines what they can say, what they cannot say yet, and what eventually breaks through despite defenses. The psychoanalytic theory is the engineering here.

2. Architecture Overview

ANIMA operates as a three-layer system with explicit interfaces between layers. The key design principle is strict separation of concerns: each layer has a single responsibility and communicates with adjacent layers through a defined interface.




Figure 1. ANIMA three-layer architecture. The LLM appears once as a renderer. All state computation is deterministic.

The critical property of this design: the LLM is interchangeable without changing the agent. Replacing GPT-4 with Claude or Llama changes the surface texture of the language — vocabulary, register, stylistic choices — but not the behavioral trajectory. The same Instruction

Object produces the same structural behavior regardless of which model renders it. This is empirically testable and is the subject of our ongoing evaluation.

3. The Psychodynamic Engine

3.1 State Variables

The engine maintains a seven-dimensional state vector updated at each conversational turn t :

$$S(t) = (E(t), \theta(t), A(t), C(t), G(t), \Pi(t), \rho(t)) \in [0, 1]^7$$

Variable	E — Eros	θ — Thanatos	A — Anxiety	C — Guilt	G — Jouissance	Π — Unc. Pressure	ρ — Fantasy Rigidity
Range	[0,1]	[0,1]	[0,1]	[0,1]	[0,1]	[0,1]	[0,1]
Psychic role	Libidinal bond	Death drive / resistance	Alarm signal	Superego pressure	Drive satisfaction	Irruption motor	Defense impermeability
Process dir.	↑ with opening	↑ with closure	↑ with $\Pi \times \sigma_{fan}$	↑ after irruption	↑ with symptom	↑ with agenda gap	↓ with genuine irruption
Histeria $t=0$	.55	.35	.40	.20	.30	.25	.75
Obsession $t=0$	.48	.45	.38	.42	.35	.30	.78
Melancholia $t=0$	.22	.65	.45	.70	.18	.35	.72
Paranoia $t=0$	.38	.55	.42	.12	.55	.38	.85

Table 1. State variable definitions, process direction, and initial values by archetype.

3.2 Update Equations

At each turn t , the engine extracts a signal vector $\sigma(t)$ from the rendered exchange and applies the following update equations ($\text{clip}(x) = \max(0, \min(1, x))$; structure-specific rate parameters $\alpha_X, \beta_X, \gamma_X$ defined per archetype):

$$\begin{aligned}
 E(t+1) &= \text{clip}[E(t) + \alpha_E \cdot \sigma_{aper} \cdot (1-\rho) - \beta_E \cdot \sigma_{cie} \cdot \theta - \gamma_E \cdot \sigma_{fan} \cdot A] \\
 \theta(t+1) &= \text{clip}[\theta(t) + \alpha_\theta \cdot \sigma_{cie} \cdot \rho - \beta_\theta \cdot \sigma_{elab} \cdot E] \\
 A(t+1) &= \text{clip}[A(t) + \alpha_A \cdot \Pi \cdot \sigma_{fan} - \beta_A \cdot \text{defense} + \gamma_A \cdot \text{irr}] \\
 C(t+1) &= \text{clip}[C(t) + \alpha_C \cdot \text{irr} \cdot k_C(\text{arch}) - \beta_C \cdot \sigma_{elab}]
 \end{aligned}$$

$$\begin{aligned}
G(t+1) &= \text{clip}[G(t) + \alpha G \cdot \sigma_{\text{int}} - \beta G \cdot \text{irr_gen} - \gamma G \cdot \sigma_{\text{elab}} \cdot E] \\
\Pi(t+1) &= \text{clip}[\Pi(t) + \alpha \Pi \cdot d_{\text{agenda}} - \beta \Pi \cdot \sigma_{\text{elab}} - \gamma \Pi \cdot \text{irr} \cdot k_{\text{irr}}] \\
\rho(t+1) &= \text{clip}[\rho(t) - \alpha \rho \cdot \text{irr_gen} - \beta \rho \cdot \sigma_{\text{elab}} \cdot \Pi(t-1) + \gamma \rho \cdot \text{defense}]
\end{aligned}$$

Equations 1–7. Psychodynamic engine update rules. All parameters are positive scalar rates; $kC(\text{arch})$ is an archetype-specific guilt multiplier.

3.3 Irruption Mechanism

The engine's central event is the irruption: the moment when unconscious pressure exceeds the agent's defensive capacity and something unplanned enters the discourse. Irruption follows a two-condition rule:

Necessary condition (deterministic): $\Pi(t) \geq \theta_{\text{irr}}(\text{arch})$
Sufficient condition (stochastic): $U \sim \text{Bernoulli}(\text{sigmoid}(k \cdot (A(t) + G(t) - \theta_{AG})))$

Irruption thresholds range from $\theta_{\text{irr}} = 0.45$ (schizophrenia: no conventional repression) to $\theta_{\text{irr}} = 0.78$ (perversion: strong disavowal containment). The type of irruption is archetype-specific: behavioral closure in hysteria, insight ritualization in obsession, somatic in phobia, direct self-accusation in melancholia, chain expansion in paranoia, signifier dissolution in schizophrenia, and genuine aperture in perversion.

4. LLM Coupling Protocol

The coupling protocol defines the interface between the psychodynamic engine and the language model. At each turn, the engine produces a structured Instruction Object that the LLM receives as its primary input. The Instruction Object specifies exactly what behavioral state must be rendered, without requiring the LLM to understand the underlying dynamics.

The Instruction Object contains seven fields:

Field	Content
<code>estado_psiquico</code>	Numerical values of all seven state variables with qualitative interpretation (e.g., <code>vinculo</code> : 'replegado', <code>presion</code> : 'alta')
<code>condicion_activa</code>	Active condition: normal defensa irrupcion post_irrupcion banal_apertura
<code>instrucciones_respuesta</code>	Behavioral pattern (ordered sequence), length, register, Gricean maxims to violate, and how
<code>contenido</code>	Top agenda item, available cultural references, relevant session history

<code>restricciones_duras</code>	Hard constraints the renderer must follow without exception
<code>marcadores_corporales</code>	Somatic channel markers: actions to be rendered as bracketed stage directions
<code>nota_para_renderer</code>	Clinical note synthesizing the psychodynamic moment — the renderer's main guide

Table 2. Instruction Object fields. The LLM receives this object and produces natural language that satisfies all constraints.

The system prompt configures the LLM as a renderer: it receives the persona's voice profile, active biographical history, and rendering rules (corporal markers in brackets, length limits, etc.). The system prompt is static per instance; the Instruction Object changes every turn. This architecture makes the LLM trivially replaceable: any model capable of following structured instructions produces the same agent.

5. Deterministic Signal Extractor (DSE)

After the LLM renders the agent's response, the engine must update $S(t+1)$. This requires computing the signal vector $\sigma(t)$ from the rendered text. An earlier version of ANIMA used a second LLM call for this classification step. The DSE eliminates that dependency entirely.

The DSE is a pure function: text in, signal vector out. No API calls, no stochasticity, no black-box inference. Four components:

5.1 Structural Pattern Analyzer → $\sigma_{cie}, \sigma_{sint}$

Each behavioral archetype has characteristic discourse patterns identifiable through regex matching. For hysteria: deflection markers, environmental comments, questions directed at the interviewer. For obsession: quantification, procedural language, ritualization. For paranoia: causal chain constructions, certainty enunciations. Each pattern has a calibrated weight; the weighted sum produces σ_{cie} and σ_{sint} .

5.2 Fantasy Object Detector → σ_{fan}

Each agent instance carries a fantasy object specification: a weighted keyword set (high-weight / medium-weight) and regular expression patterns derived from the instance's fantasy formulation. The DSE scores the rendered response against this specification. If the weighted score exceeds a calibrated threshold (default: 0.42), $\sigma_{fan} = 1$.

5.3 Agenda Proximity Calculator → d_{agenda}

The engine maintains an internal agenda — a priority-ordered list of items the agent 'needs' to communicate. The distance between what the agent said and what the top-priority agenda item contains is computed as:

$$d_{\text{agenda}} = \max(0, 1 - \text{cosine_sim}(\text{response_bow}, \text{agenda_top_bow}) \times 3.5)$$

Where bow denotes bag-of-words representation after stopword removal. When d_{agenda} is high, the agent is talking about something other than what presses most — and Π accumulates.

5.4 Elaboration Detector → σ_{elab}

Genuine elaboration — a subject making connections across previously separated material — has a quantifiable signature: semantic novelty relative to recent turns. The DSE computes cosine similarity between the current response and the preceding six responses. High novelty (low max similarity) indicates elaboration. Connecting language patterns (e.g., 'I realize', 'it's always the same') provide additional signal boosts.

5.5 Aperture Analyzer → σ_{aper}

Discursive aperture is measured through three syntactic features: first-person pronoun ratio (self-reference vs. other-reference), absence of interviewer-directed questions (deflection markers), and sentence density. These combine into σ_{aper} without semantic inference.

DSE Property:

Every σ value computed by the DSE is fully traceable: the user can inspect which patterns fired, which keywords matched, what the cosine similarity was, and how each contributed to the final value. This auditability is absent from any LLM-based classification approach and is essential for scientific reproducibility.

6. Behavioral Archetypes

ANIMA currently implements seven behavioral archetypes drawn from psychoanalytic structural theory. Each archetype is a parameter bundle: initial state values, rate multipliers, irruption threshold, defense mechanism, Gricean violation profile, irruption type, and cultural world specification. The archetypes are not character templates — they are attractor configurations in the $S(t)$ state space.

Archetype	Defense mechanism	Irruption type	Discourse signature
Hysteria	Repression	<i>Behavioral closure</i>	Deflection, other-reference, self-interruption before nucleus
Obsession	Affect isolation	<i>Ritualized (insight → task)</i>	Quantification, procedural language, excess information
Phobia	Inhibition / avoidance	<i>Radical verbal</i>	Circumlocution around object, somatic surge, 'falling nowhere'

Melancholia	Introjection	<i>Direct self-accusation</i>	Minimal response, flat affect, self-as-destroyer formulations
Paranoia	Projection	<i>Chain inclusion of therapist</i>	Causal certainty, evidence accumulation, no contingency
Schizophrenia	Foreclosure (structural)	<i>Signifier dissolution</i>	Functional neologisms, tangential discourse, absent anxiety
Perversion	Disavowal	<i>Genuine aperture</i>	Simultaneous assertion/denial, irony, jouissance around transgression

Table 3. Seven behavioral archetypes. Each archetype is a full parameter bundle, not a character description.

7. Preliminary Evidence and Validation Protocol

7.1 Demonstration Sessions

We have run demonstration sessions for all seven archetypes under a student-of-psychology interaction protocol. Sessions consist of 4-6 exchanges with no explicit behavioral prompting — the interviewer asks open questions; the archetype responds according to its structural parameters. Key findings:

- Irruption events occur without scripting, as emergent consequences of Π exceeding θ_{irr} . The form of irruption is archetype-specific: hysteria closes behaviorally, obsession ritualizes, paranoia includes the interviewer.
- ρ (fantasy rigidity) shows consistent decrease across sessions for neurotic archetypes — a computational analogue of therapeutic process. Paranoia shows minimal ρ decay, consistent with clinical theory.
- σ_{fan} activation correctly tracks theoretically relevant questions: the hysteria instance's fantasy object activates on choice/abandonment questions, the paranoia instance's on contingency challenges.

7.2 Blind Clinical Evaluation Protocol

We have designed and are executing a Level-1 validation study: blind classification of ANIMA transcripts by clinical evaluators with psychoanalytic training. Protocol summary:

Parameter	Value
Corpus	7 transcripts (one per archetype), 4 exchanges each, presented in randomized order without structure labels
Evaluators	N = 10-15; inclusion criteria: practicing psychoanalytic clinician or advanced trainee

Task	Classify structure (7 options), rate confidence (1-5), describe markers, indicate differential if confidence ≤ 3
Primary measure	Classification accuracy vs. chance ($1/7 \approx 14.3\%$); binomial exact test, $p < 0.05$
Secondary measure	Inter-rater kappa (target: $\kappa \geq 0.30$); confusion matrix for development guidance
Success threshold	Accuracy $\geq 40\%$ ($\approx 3 \times$ chance) AND $\kappa \geq 0.30$

A positive result provides publishable empirical evidence of structural distinguishability. A negative result provides a confusion matrix that precisely identifies which archetypes require further engine development — a result equally valuable for the research program.

8. Applications

8.1 Agent Behavioral Evaluation

ANIMA's DSE is a prototype methodology for auditable behavioral evaluation. Current agent evaluation frameworks assess surface properties (task completion, factual accuracy). Structural consistency over time — whether an agent's behavioral trajectory is coherent, whether it maintains appropriate internal pressure, whether its responses are generated from an actual internal state — is not evaluated. The DSE's deterministic signal extraction provides a framework for this evaluation that requires no additional LLM calls and produces traceable, reproducible results.

8.2 Structured Synthetic Data Generation

Training data for conversational agents is predominantly flat: behaviorally undifferentiated, without internal coherence, without the dynamics of accumulating pressure and release. ANIMA can generate structurally diverse conversational corpora — seven distinct behavioral organizations, each with internally consistent trajectories across sessions. Synthetic data generated through ANIMA carries behavioral structure that current synthetic data generation cannot provide.

8.3 Clinical and Training Applications

Clinical training in psychoanalytic or psychodynamic psychotherapy requires exposure to diverse structural presentations — not always available in supervised practice. ANIMA can generate training cases with calibrated structural properties, including controlled irruption events at known turn numbers, for supervised learning environments. The blind evaluation protocol we have designed functions simultaneously as a validation study and as a training assessment tool.

8.4 Persistent NPC and Character Systems

Game and interactive narrative systems require characters that behave consistently over long interactions and accumulate history. Current LLM-based NPCs lose behavioral coherence beyond a few turns. ANIMA's architecture provides a principled solution: persistent internal state that the LLM renders, rather than state maintained by the LLM itself.

8.5 Market Simulation and Behavioral Economics (Future Direction)

Current agent-based market simulation models face a structural limitation: their agents either follow simple behavioral rules or receive a one-time personality prompt from an LLM. Neither approach can model the accumulation of internal pressure across repeated exposures, the evolution of a subject's position over time, or the differential response to messages based on the agent's underlying subjective organization. The result is simulations that answer demographic questions rather than psychological ones.

ANIMA's psychodynamic engine is a natural candidate for richer simulation. Instead of asking 'what percentage will buy?', a market built on ANIMA's architecture could address structurally different questions:

- Which message generates initial adhesion across different subjective organizations — and which produces immediate resistance?
- Which profiles change position after repeated exposure, and which develop resistance after initial interest?
- How does acceptance evolve after multiple campaign interactions over time?
- Which segments become promoters, and what internal state profile predicts that transition?

In a streaming service launch scenario, for example, a synthetic market of subjects with diverse psychodynamic state distributions could model how different campaign messages interact with different internal economies — not as statistical correlations but as structural responses derived from the agent's state. A subject with high Π (unconscious pressure) and low ρ (fantasy rigidity) responds differently to a scarcity message than a subject with high G (jouissance) and high ρ . That difference is not a personality trait — it is a dynamic property of the agent's internal state at the moment of exposure.

This distinguishes the approach from current digital twin and consumer simulation platforms, which model behavior as a function of demographic attributes or static preference rules. ANIMA models behavior as a function of a persistent, evolving internal state — which means the same message, delivered to the same agent at different moments in a trajectory, can produce different outcomes.

Current status:

This application requires infrastructure development beyond the current prototype — specifically, a backend capable of running thousands of concurrent instances with inter-agent social influence modeling. We are currently focused on validating the core architecture in the clinical domain, which provides the empirical foundation for market

simulation applications. We are open to exploring this direction with partners who have simulation infrastructure and interest in richer behavioral modeling.

8.6 Political Communication and Discourse Research (Future Direction)

Political communication research faces a methodological constraint: studying how different message types interact with different psychological profiles requires either large-scale human subject studies — expensive, slow, and ethically constrained — or simplified computational models that lack psychological depth. ANIMA offers a third option: a simulator of psychological responses to discourse, not a predictor of individual decisions.

The distinction is precise and consequential. ANIMA can model how a given type of subjective organization responds to a given type of message — which profiles show greater resistance, which show greater openness, which tend toward polarization. It cannot and should not be used to identify real individuals with those profiles and target them accordingly. The unit of analysis is the structural type, not the person.

Legitimate experimental uses under this framing include:

- Exposing the same political message to thousands of simulated subjective states and observing which profiles show greater resistance, greater adhesion, or greater polarization — as a tool for understanding structural vulnerabilities in public discourse, not for exploiting them.
- Comparing how different communication styles generate different internal responses by subjective organization — studying what makes a message land as persuasion versus imposition versus threat, as a function of the listener's internal state at the moment of exposure.
- Studying which types of discourse favor dialogue versus confrontation — identifying structural features of communication that increase or decrease the probability of genuine elaboration in the interlocutor.

These applications are closer to computational political psychology than to political strategy. The research question is structural: what does a given type of discourse do to a given type of internal organization? Not: how do we reach and persuade this specific population?

Ethical boundaries and acceptable use

The political communication application of ANIMA raises risks that must be named explicitly. We identify three uses that fall outside the acceptable scope of this architecture:

- Psychological microtargeting: using ANIMA-derived profiles to identify and target real individuals or populations based on inferred internal state. This application is excluded from any collaboration or licensing arrangement.
- Influence campaign design: using ANIMA simulations to optimize persuasion strategies for deployment against real populations. The architecture is designed for understanding structural responses, not for engineering them at scale.
- Predictive profiling: presenting ANIMA outputs as predictions of individual behavior or decision. The system models structural tendencies, not individual trajectories.

The distinction between a simulator of structural psychological responses and a tool for behavioral manipulation is architectural — it depends on whether the output is used to understand patterns or to target persons. We commit to this distinction as a condition of any collaboration involving this application domain, and we welcome institutional review for research uses in political communication.

8.7 Social Influence Layer (Prototype, Validated)

A fourth layer extends the architecture from a single subject in dialogue to a population of subjects embedded in a social graph. This layer does not modify the psychodynamic engine described in Section 3 — it modulates the signal vector $\sigma(t)$ that feeds it. This preserves the core architectural principle: the engine remains a pure, auditable function; social context enters only as an additional signal contribution, computed by a separate deterministic module.

Each subject in the population carries the same seven-variable state $S(t)$ as in Section 3, plus a scalar position $p(t) \in [0,1]$ representing adherence to an external message (e.g., a campaign, a policy announcement, a product launch). At each simulation tick, five factors modulate how strongly a message shifts $p(t)$ for a given subject:

- Authority — weight of the message source's perceived authority.
- Consensus — perceived proportion of the population already adhering to the message.
- Group influence — mean position of the subject's direct neighbors in the social graph.
- Source trust — accumulated reliability of the specific message source for this subject.
- Reputational cost — social cost of diverging from the perceived majority position.

Critically, each archetype has a distinct permeability profile across these five factors — derived from the same structural logic that governs the engine's defense mechanisms. Hysteria, organized around recognition by the Other, shows high permeability to consensus and group influence. Obsession, organized around procedural control, responds to authority but is largely indifferent to consensus. Paranoia inverts the sign: authority and consensus actively repel rather than attract, consistent with a structure organized around the Other's malevolent intent. The same campaign, delivered with the same strength, produces structurally differentiated population-level outcomes — not because of randomness, but because of the receiving structure's characteristic relation to the social field.

Validation

Before implementation as an interactive prototype, the social layer's dynamics were validated computationally across four tests, run deterministically with seeded random generation for reproducibility:

Test	Result
------	--------

Stability (100 ticks, 3 campaign strengths)	Converges in all conditions; no divergence or explosion
Structural differentiation	Range of 0.15 between most- and least-adherent archetypes; paranoia correctly resists more than hysteria adheres
No runaway polarization	0% of nodes reach extreme positions (<0.05 or >0.95) even under maximum campaign strength
Determinism	Identical seed produces byte-identical simulation trace

A working prototype implements this layer over a 36-node small-world graph (ring lattice with random long-range rewiring, following Watts-Strogatz topology), with a subset of nodes designated as high-authority influencers. The prototype renders the propagation of a campaign in real time, with per-archetype adherence tracked as the simulation runs.

Relation to Section 8.5

This layer is the mechanism that makes the market simulation application (Section 8.5) concrete rather than aspirational: it is what allows a population of psychodynamically distinct synthetic subjects to respond differentially, and propagate that response through a social structure, to a single piece of communication — which is the operational definition of the questions posed in that section (which message generates adhesion, which produces resistance, how acceptance evolves).

9. Current Status and Development Roadmap

Phase	Status
Psychodynamic engine (7 variables, update equations, irruption mechanism)	Complete — implemented in JavaScript, running in browser
LLM coupling protocol (Instruction Object, Pattern Library, system prompt builder)	Complete — tested with Claude claude-sonnet-4-6
Deterministic Signal Extractor v1	Complete — structural regex + keyword matching + cosine similarity, zero LLM calls
Behavioral archetypes (7)	Complete — parameter bundles with cultural world calibration for Argentine urban context
Blind clinical evaluation — corpus	Complete — 7 transcripts, randomized, formatted for evaluators

Blind clinical evaluation — execution	In progress — evaluator recruitment via clinical collaborator network
Inter-session memory P(n) and structural plasticity	Specified, not yet implemented
DSE v2 — lightweight embeddings for semantic proximity	Specified, not yet implemented
Social influence layer (5-factor modulation, 36-node graph prototype)	Validated — stability, differentiation, no-runaway, determinism all confirmed
Multi-instance interaction (transference as co-construction)	Research stage

The current implementation runs as a single-file React application with API calls to Claude claude-sonnet-4-6 as the renderer. The engine, DSE, and coupling protocol are implemented in pure JavaScript with no external dependencies beyond the React framework. The system is functional and demonstrable today.

10. Collaboration Opportunities

We are seeking research collaboration or sponsored development engagement with AI labs interested in the following problem spaces:

- Behavioral consistency evaluation: applying DSE-derived methodology to evaluate structural coherence in existing agent systems.
- Structured synthetic data: using ANIMA to generate behaviorally differentiated conversational corpora for agent training.
- Architecture research: extending the decoupling principle — deterministic state engine + LLM renderer — to other behavioral domains.

We are open to research collaboration agreements (with IP retained by project), sponsored research engagements, and technical consulting arrangements. We are not seeking acquisition at this stage.

Contact:

Project ANIMA — available for technical exchange and demonstration. Interested parties may request access to the full technical specification, live demonstration of the running system, and the validation protocol materials.

References

- [1] Freud, S. (1923). The Ego and the Id. Standard Edition, Vol. 19.
- [2] Freud, S. (1926). Inhibitions, Symptoms and Anxiety. Standard Edition, Vol. 20.
- [3] Lacan, J. (1966). *Écrits*. Seuil.
- [4] Lacan, J. (1972). On Psychoanalytic Discourse. Milan Lecture.
- [5] Miller, J.-A. (1966). Suture: Elements of the Logic of the Signifier. *Screen*, 18(4).
- [6] Recalcati, M. (2012). *Clinica del vuoto*. Franco Angeli.
- [7] Grice, H.P. (1975). Logic and Conversation. In P. Cole & J. Morgan (Eds.), *Syntax and Semantics*. Academic Press.
- [8] Damasio, A. (1994). *Descartes' Error: Emotion, Reason, and the Human Brain*. Putnam.
- [9] Friston, K. (2010). The free-energy principle: a unified brain theory? *Nature Reviews Neuroscience*, 11(2), 127-138.
- [10] Landis, J.R. & Koch, G.G. (1977). The measurement of observer agreement for categorical data. *Biometrics*, 33(1), 159-174.
- [11] Watts, D.J. & Strogatz, S.H. (1998). Collective dynamics of 'small-world' networks. *Nature*, 393(6684), 440-442.